



MapReduce Project

Matrix-Multiplication

Hardik Shah

Kevin Shah

Ruturaj Nene



Overview

- Matrix Multiplication Sparse H-V
- PageRank Sparse H-V
- Matrix Multiplication Dense B-B
- Analysis and Experiments




Matrix Multiplication - Sparse H-V partitioning

- Partition left matrix horizontally (H)
- No need to partition right matrix ($N \times 1$ Vector)
- Duplicate right vector H times
- Each reducer receives group of rows from left matrix and whole right vector, multiplies each row with the column vector



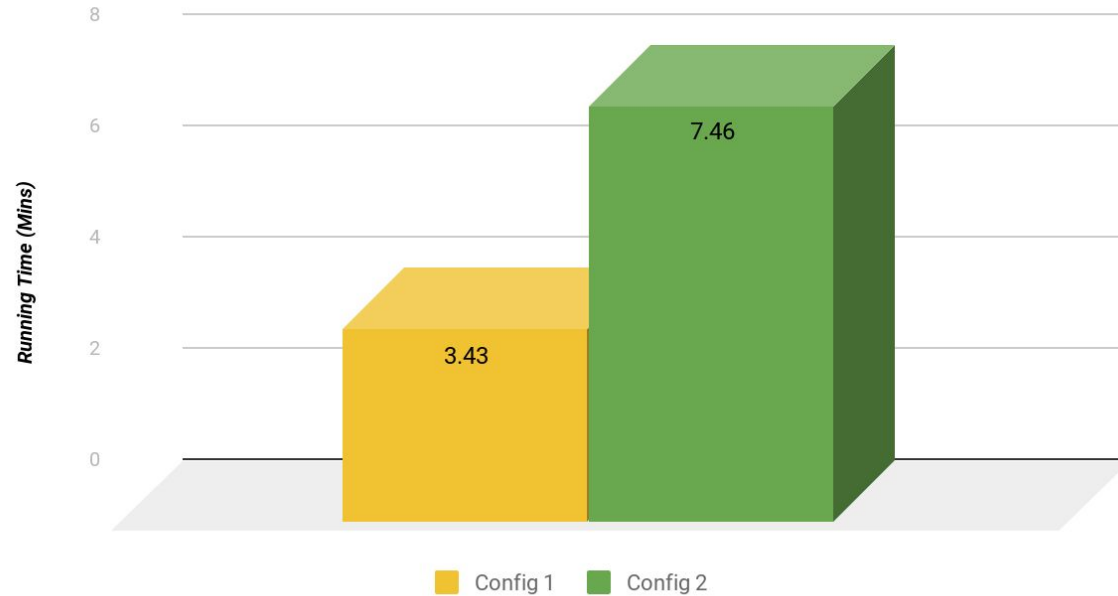
Speedup



Config 1- Matrix size:
30k x 30k and 30k x 1, 4
workers, 20
Partitions(H)

Config 2- Matrix size:
30k x 30k and 30k x 1, 8
workers, 20
Partitions(H)

Scalability




Config 1- Matrix size:
20k x 20k and 20k x 1, 4
workers, 20
Partitions(H)

Config 2- Matrix size:
30k x 30k and 30k x 1, 4
workers, 20
Partitions(H)




Partition Granularity



Config 1- Matrix size:
30k x 30k and 30k x 1, 4
workers, 20
Partitions(H)

Config 2- Matrix size:
30k x 30k and 30k x 1, 4
workers, 20k
Partitions(H)


Pagerank analysis



Config 1- Synthetic graph from HW4, $k=500$, 10 workers

Config 2- Matrix size: $250k \times 250k$ and $250k \times 1$, 10 workers, 25 Partitions(H)

Matrix Multiplication (Dense B-B partitioning)





Matrix Multiplication (Dense B-B partitioning)

- JOB 1: Divide left and right matrix into block partitions ($H1 \times V1$, $V1 \times V2$) and computes partial sums of certain matrix products
 - Mappers:
 - A_{ij} -> key = (I,J,K) for all K Partitions, value = ("A", i, j, A_{ij})
 - B_{jk} -> key = (I,J,K) for all I Partitions, value = ("B", j, k, B_{jk}), where i,j,k are part of partitions I,J,K respectively
 - Reducers:
 - For key = (I, J, K) compute $x_{ik} = \sum A_{ij} * B_{jk}$ For all j in J, which is the partial sum for matrix product C_{ik} corresponding to a partition J




Matrix Multiplication (Dense B-B partitioning)

- JOB 2: Sum the partial sums and emit the matrix product
 - Mappers:
 - $x_{ijk} \rightarrow \text{key} = (i), \text{value} = (k, x_{ijk})$
 - Reducers:
 - For key = (i), accumulate partial sums for all k in a HashMap over all J Partitions
 - Emit output $C_{ik} = \sum x_{ijk}$ for all J Partitions, for all k in the HashMap



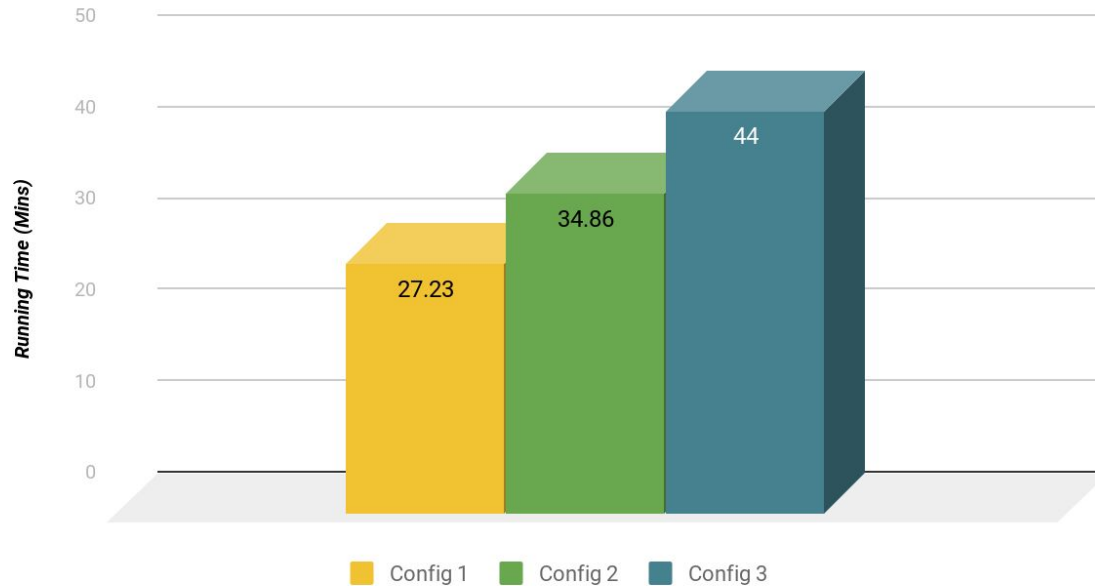
Speedup



Config 1- Matrix size: 6k x 6k and 6k x 6k, 5 workers, H1, V1, V2 = 10

Config 2- Matrix size: 6k x 6k and 6k x 6k, 10 workers, 10, H1, V1, V2 = 10

Scalability

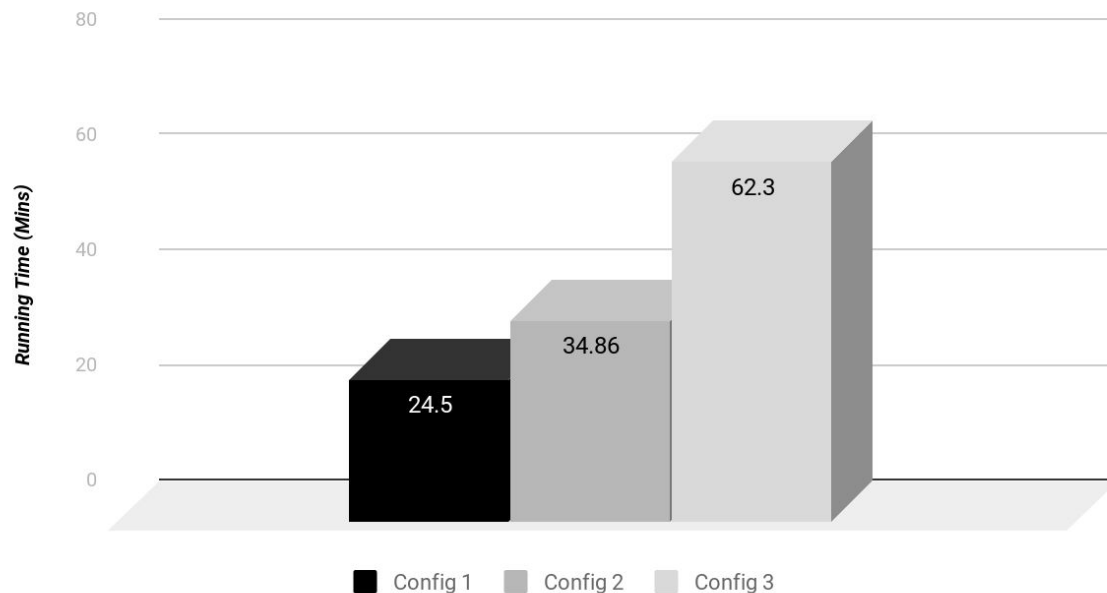


Config 1- Matrix size: 5k x 5k and 5k x 5k, 10 workers, H1, V1, V2 = 10

Config 2- Matrix size: 6k x 6k and 6k x 6k, 10 workers, H1, V1, V2 = 10

Config 3- Matrix size: 7k x 7k and 7k x 7k, 10 workers, H1, V1, V2 = 10

Partition Granularity



Config 1- Matrix size: 6k x 6k and 6k x 6k, 10 workers, H1, V1, V2 = 5

Config 2- Matrix size: 6k x 6k and 6k x 6k, 10 workers, H1, V1, V2 = 10

Config 3- Matrix size: 6k x 6k and 6k x 6k, 10 workers, H1, V1, V2 = 20



Conclusion

- Achieved good scalability and speedup on Sparse H-V and Dense B-B matrix multiplication
- Proved concepts taught in class like increasing partition granularity can lead to worse performance because of higher duplication
- Dense B-B can be used for any application that requires matrix multiplication. Eg: Pagerank